

Build Log : Application Prep Agent

Platform

Claude Project ("Job Application Tailoring"), free tier. Custom instructions pasted from the FL-06 spec. Project knowledge: master CV, cover letter template, proof statement. Live connector: Google Drive (MCP), enabled at the Project level.

Narrowest end-to-end version, working

Core loop: human pastes one real job posting → agent checks Drive for source documents → extracts requirements from the posting → checks for hard eligibility gaps → drafts, critiques, and revises a tailored CV/cover letter, or names the gap and recommends skip → final recommendation, no auto-send. First full run completed against a real posting (Turing) before any other feature was added, per the brief's own sequencing.

What broke, what changed

This is the real iteration, in the order it happened and not a cleaned-up summary.

What happened	What broke / what I found	What I changed
Gather step, first real run	NotebookLM isn't callable as a tool inside this platform, no direct integration available in-chat	Deviation from FL-06 spec: extraction done directly against fetched posting text instead of via NotebookLM. Documented, not hidden.
Fetching a LinkedIn posting (Turing)	Direct URL fetch fails, LinkedIn blocks automated access even by exact link	Gather step requires the human to paste posting text manually for gated sites. Not fixable inside this build.
Checking a claimed skill (TensorFlow / gesture project)	GitHub's repository list page returns ROBOTS_DISALLOWED; a specific unlisted repo also didn't surface via search	Refused to include the claim until it could be verified from a real source — later confirmed directly once the user's updated CV included it with real metrics (96.7% test accuracy)
CV credential check (Kuda run)	Uploaded CV said "MSc, Data Mining & Knowledge"; user later stated "MSc Computer Science"; a real discrepancy across their own documents	Guardrail fired as designed: stopped and asked rather than picking one silently. Resolved once user confirmed and updated the master CV in Drive.
Blend360 posting	User initially disputed a "required" reading of the Databricks line	Re-checked the exact source text rather than deferring; confirmed Databricks and SQL Server were both listed under the posting's actual Required section, distinct from a separate Databricks-certification perk mentioned elsewhere. Held the original assessment with the exact quote shown.
Tracker feature request	Google Drive connector has create_file but no update_file, an existing Sheet's content can't be edited, only replaced	Built as a single recreated snapshot rather than an appendable log. Real limitation, not a design choice, noted for a future version.

Real decisions made across runs

Seven postings run through the pipeline so far. Three genuine skips, not just happy-path drafts, evidence the eligibility guardrail actually fires on different real inputs.

Posting	Outcome	Why
Turing, Analytical AI Trainer	Drafted, ready to send	Strong fit after critique/revise; added truthful keyword coverage (research, collaboration, LLM comfort)
Kuda, Data Analyst, Credit	Skipped	Hard gap: 3+ yrs fintech-specific experience required, not held; not stretched
Burjlinebuilders, Data Quality Analyst	Drafted, ready to send	Strong fit; re-run after CV update added a verified 1M-record SQL project as new evidence
Outreachy Internship	Blocked, not drafted	No extractable posting content, initial applications weren't even open yet at time of check
Blend360, Data Engineer Lead	Skipped	Hard gaps: SQL Server + Databricks required, Lead-level experience required, neither held

Valcet, Analytics Engineer	Skipped	Hard gap: posting states 10+ years' experience as a Must-have
Qureos, Data Quality Analyst	Drafted, ready to send	Strong fit after posting check and No hard eligibility gap

Deviations from the FL-06 spec, summarized

- NotebookLM replaced with direct extraction, not available as a callable tool in this environment.
- Auto-apply was already out of scope by design in FL-06; unchanged here, every run ends in a human decision, not a submission.
- Tracker (job application log) added mid-build, beyond the original spec, a genuine scope addition, not a cut. Implemented as a Google Sheet via the Drive connector's create_file, recreated per update rather than appended, due to a real tool limitation (no update_file).

Supporting live artifact

Job Application Tracker (real Google Sheet, created by the agent via a live Drive write action, not just a read): [view sheet](#)

Recorded run

Qureos, Data Quality Analyst, run end-to-end inside the actual Claude Project (not rehearsed here), screen-recorded raw and unedited, ~3 minutes, matching the same core loop demonstrated across every rehearsal run above.